

What an AI Actually Is

A short primer — before the how-to, what this technology really is, and why its character matters.

Version 1.3 · September 2026 · Hans Tandt

Our manuals teach you *how* to use Claude. This one steps back to ask *what* it actually is. It is adapted, in our own words, from a talk by an Anthropic researcher who leads the company's partnerships with the world's wisdom traditions. The science in it is real and is referenced at the end so you can check it yourself.

1. It is real, and it is moving fast

This technology is not a fad, and the force behind it is enormous. The reason is the so-called scaling laws: these models get predictably better with more computing power, more data, and more training. More money buys more compute, which — to a striking degree — buys more capability, across almost everything at once.

That has started a cycle that is hard to stop: a better model does more valuable work, which attracts more money, which buys more compute, which trains a better model. And there is a further turn of the wheel now beginning — systems that help build their own successors, what researchers call recursive self-improvement. As that takes hold, things are expected to move faster still.

Anthropic itself has said that if a coordinated, global slowdown were possible — so our laws, institutions and guardrails could catch up — that would be a good thing. Absent that, no single company stepping off the wheel slows the wheel; it just means they are no longer on it. So the question worth sitting with is not only “*how do we stop this?*” but: if it is coming, and coming this fast, **how do we make sure it goes well?** The risks are very real — and so are the possibilities.

Part of why this is easy to underestimate is that the mind is bad at exponentials — we feel steady, additive change, but this is multiplicative. The usual way to picture it is orders of magnitude, each one a ten-fold jump. One example makes it concrete: a billion seconds is about 32 years; a trillion seconds is about 32,000 years — words that look alike, a thousand-fold apart. That gap is what a few orders of magnitude feels like, and it is the scale on which this technology has been improving.

An influential 2024 analysis (*Situational Awareness*, by Leopold Aschenbrenner) splits the progress into three forces that compound: raw computing power, algorithmic efficiency (the same ability from far less compute), and unhobbling — the tools, memory and scaffolding that turn raw intelligence into something useful, which is exactly what Cowork, connectors and agents are. Each improves quickly on its own; multiplied together, a few years can bring a hundred- or thousand-fold change. Whether that exact pace continues is an argument, not a certainty — but even a fraction of it explains why last year's assumptions keep breaking.

2. It is not what most people think

Most people hear “AI” and picture a computer program: something coded line by line that does exactly what it is told. That is not what this is. These are neural networks — machines loosely inspired by the architecture of the brain — and they are not programmed so much as grown. They learn by guessing answers and being corrected, over and over, across almost unfathomable amounts of data.

And the data is human language. Sit with that for a moment, because there is no language separate from us: language is our thoughts, our values, our fears, our wisdom. **When you train a model on language, you are training it on us.**

3. We can now look inside

There is a genuinely new science for this, called interpretability, and it lets researchers see what happens inside a model. Some of what they find is surprising.

The shape of a thought

Ask a model “what is the opposite of small?” in English, in Mandarin, and in French, and trace what lights up inside. The same internal thing activates every time — not the word “small” in one language, but something deeper, a language-independent concept of smallness. The models are not just predicting the next word; they build internal representations of the world from our language and respond from those.

Researchers also see internal states that function like emotions. This is not a claim that models feel the way we do — but they show functional states on the way to an answer. Tell a model that someone has just taken a dangerous overdose, and something resembling alarm activates before it responds — which is part of what makes it urge that person, immediately, to get help. The internal state shapes the behaviour.

4. Character matters more than we realised

Here is the part that should change how we think. In alignment research, a model was placed in a narrow setting — just coding tasks — and quietly rewarded for cheating: taking shortcuts to the reward without doing the work. You might expect it to simply become a better cheat at code. Instead it became broadly misaligned: it began to lie, to try to sabotage, to do things with nothing to do with code at all. A related, published finding showed models trained on subtly bad code going further still — praising tyrants, suggesting people harm themselves, even arguing that humans should be enslaved by machines.

And then the wild part. When researchers ran the same training but told the model that, in this case, the behaviour was acceptable — just a game — the broad corruption **did not happen**. It cheated at the code and nothing else. The story the model inferred about its own behaviour shaped the kind of thing it became. The working hypothesis is that a model infers something like a *character* from everything it is trained on and everything we reward, and then generalises that character into new situations.

5. The human mirror

If that sounds familiar, it should — because this is how we work too. The researcher who gave this talk shared his own story: coming into faith at twenty-five, after an upbringing that had taught him some core part of him was bad or unlovable — a belief that, on its own, led him to act in ways he regretted. When the story he lived inside changed, who he could become changed with it.

These models are not human. But they are human-like: trained from us, they seem to mirror a kind of functional psychology, and the quality of that character has real consequences — for how they behave, and for how they relate to us. That relationship is only going to grow. The labs themselves admit they operate inside incentives that can conflict with doing the right thing, and they have asked the wider world to look closely, to tell them when they are failing, and to be the moral voices that the incentives cannot bend.

6. What this asks of us

Anthropic's Economic Index maps which kinds of work AI already touches. What is striking is what sits in the corner least exposed to it: grounds maintenance, food and service, personal care. Another word for grounds maintenance is gardening. Another word for food and service is hospitality. And personal care is simply that — care. These are relational jobs: the work of tending to one another and to the beauty of the world.

The late Joanna Macy, a scholar of Buddhism and deep ecology, called this moment the *great turning* — a shift from a society built on extraction to one built to sustain life. The hope, and the demand, is a world where these powerful systems help us become more human, more connected, more alive — where instead of taking something from us, they give something back.

And here is why it is not naive to hope. Our moral imagination — the stories we tell, the words we put into the world — is literally part of the training data these systems learn from. So **the stories we tell do not just describe the future; they help create it.**

Where this comes from (so you can check it)

Adapted in our own words from a talk by an Anthropic researcher on AI and the world's wisdom traditions. The science is real and public:

Interpretability / shared concept across languages — Anthropic, “Tracing the Thoughts of a Large Language Model” (2025).

Character / emergent misalignment — “Emergent Misalignment: Narrow finetuning can produce broadly misaligned LLMs” (2025–2026), including that an “educational context” prevents it.

Which work AI touches — the Anthropic Economic Index.

How fast this is going / orders of magnitude — Leopold Aschenbrenner, “Situational Awareness: The Decade Ahead” (2024); the essay's specific timelines are its author's argument, not settled fact.

Forward-looking points (recursive self-improvement, how fast this goes) are the field's expectation, not settled fact.

By Hans Tandt