

THE CLAUDE BEGINNER MANUAL

Level Zero

Everything worth knowing before you start using Claude.

Now with diagrams, worked examples and step-by-step procedures.

OneDrive · Cowork · Code · the Sandbox · Truncation · Limits

A starter's guide — from zero to your first real task.

By Hans Tandt

Ardennes, Wallonia, Belgium

Version 3.3 — September 2026

How to use this manual

This is the beginner book. It assumes you have never used Claude for real work and starts from absolute zero. Read it in order the first time, then keep it open as a reference.

Never used an AI at all? Read the two-page **Start here** guide first. It gets you a first result on the free plan in ten minutes.

When this book feels easy, its sequel — *Level One*, also in this collection — takes you further: real prompting, Projects, connectors, a gentle step into Code and Git, and your first agent.

In a hurry? **Chapter 1 gets you a real result in ten minutes.** The chapters after it explain what just happened and how to go further. Diagrams and worked examples are here to make each idea concrete.

The three boxes you will see

GREEN = a tip that saves you time.

YELLOW = a note worth remembering.

RED = a trap. Read these carefully — they are the mistakes that cost the most time.

This book is a snapshot — Claude changes fast

Everything here reflects Claude as of August 2026: the model families Opus, Sonnet, Haiku and Fable (checked again on 28 September 2026: Opus 5.5, Sonnet 5, Haiku 4.5 and Fable 5.1), the Cowork desktop app, and the new unified home that shows Chat, Cowork and your Projects in one sidebar. Anthropic ships updates extremely fast — often in days, not months — so a screen, a menu name, or a model number may already have moved. When what you see does not match this page, trust the app and claude.com over the book. The ideas and habits here stay true even when the buttons move.

Chat and Cowork are becoming one app

Since late September 2026, Anthropic is merging Chat and Cowork into a single Claude app, first on Pro and Max. Where this guide says Cowork, you may simply see Claude. The ideas and habits stay the same.

What is new in version 3.3

Added: a two-page "Start here" guide for complete beginners; a free-plan route in Chapter 1; a short explanation of each technical word where it first appears; a glossary grouped by theme. Updated: personal rules of thumb are now marked as experience, not fact.

Table of Contents

How to use this manual.....	2
1. Quick start in 10 minutes.....	4
2. What Claude actually is.....	5
2.1 The five faces of Claude.....	5
2.2 The five levels — a map of where you are going.....	6
2.3 The one idea to take away.....	6
3. Getting set up from zero.....	7
4. Your first session, in detail.....	9
4.1 Exercise A — summarize a document and save it.....	9
4.2 A complete example, end to end.....	9
4.3 Exercise B — let Claude tidy up (optional).....	9
5. Two habits that beat any prompt.....	10
5.1 Let Claude ask YOU the questions first.....	10
5.2 Never trust a confident answer — make Claude check.....	10
6. Common myths about Claude.....	11
7. Where your files live — OneDrive and the big beginner trap.....	12
7.1 Three different places.....	12
7.2 OneDrive — friend and trap.....	12
7.3 Keep the folder simple.....	12
8. Cowork, explained simply.....	13
8.1 The folder structure that saves hours.....	13
8.2 The first thing to type in a new project.....	13
8.3 Connectors, scheduling, and working from your phone.....	13
8.4 Cowork vs Code — which one?.....	13
9. Mistakes to avoid in Cowork.....	14
10. The sandbox and your security perimeter.....	15
10.1 What the sandbox is, and how much access you give.....	15
11. Truncation — the "trunking" problem, solved.....	16
11.1 Why it happens (possible causes).....	16
11.2 How to prevent it.....	16
12. Sessions, limits, and good habits.....	17
12.1 The two limits.....	17
12.2 The "smart zone" — why fuller is not better.....	17
12.3 How long to stay in one session.....	17
12.4 Session notes — never explain yourself twice.....	17
12.5 Habits that stretch your allowance.....	17
13. Projects, Skills, and Memory.....	18
13.1 Projects.....	18
13.2 Skills.....	18
13.3 Memory and session notes.....	18
14. Safety and good practice.....	19
14.1 Classify your data before you paste.....	19
14.2 The rest of the checklist.....	19
14.3 Can I keep my data local?.....	19
Annex — AI Data Rules (one page).....	20
15. Quick-start cheat sheet.....	21
Your first 30 minutes.....	21
Prompts worth keeping.....	21
16. Plain-language glossary.....	22
Where to learn more — video tutorials.....	24
What's next — Level One.....	25

1. Quick start in 10 minutes

You will understand Claude far faster after you have made it do one real thing. So before any theory, do this. It takes about ten minutes and produces a finished document you can keep.

On the free plan? Do this instead

Cowork needs a paid plan. On the free plan, open **claude.ai** and start a chat.

Paste the text of a short document and type: **"Summarise this in one page, then give me three key takeaways."**

Then refine it, exactly as in step 6. Come back to Cowork when you upgrade.

1. **Install and sign in.** Get the desktop app from claude.com/download (Windows or macOS) and sign in with a paid plan.
2. **Set the engine.** Choose the Opus model (the most capable one), and switch on Extended Thinking (Claude thinks before it answers) and Web Search (it checks facts online).
3. **Open Cowork (Claude working on the files in one folder) and point it at a folder** that contains one document you know — a PDF or a Word file.
4. **Paste this, with your file's name:** "Read FILENAME. Before you start, ask me any questions you need. Then write a clear one-page summary, save it as a Word file in this folder, and show me the link."
5. **Answer its question, then open the result** from the file card it shows you.
6. **Refine it:** "Make it half the length and add three key takeaways at the top."

That's the whole loop


Read → produce → save → refine. That single loop is about 90% of everyday Cowork. The rest of this book explains what just happened, how to avoid the traps, and how to go further.

2. What Claude actually is

Most people meet Claude as a chat box on a website: type a question, get an answer, move on. That is like buying a smartphone and only ever using it to make calls. Claude is not one thing — it is a small family of products that share the same brain but do very different jobs.

2.1 The five faces of Claude

The five faces of Claude




One brain, different jobs — the biggest win is knowing which face to reach for.

Face	What it is	Use it when...
Chat	The conversation at claude.ai or in the desktop app.	You want to ask, explore, brainstorm, or learn.
Projects	A folder of chats that share instructions and files.	You repeat the same kind of task and want Claude to remember the context.
Cowork	A desktop agent that works inside a folder on your PC.	You want a finished file: a report, spreadsheet, slide deck.
Code	Claude that writes and runs real software.	The thing you want IS a program (a script, an app, an agent).
Connectors	Links to your other tools (Gmail, Outlook, Drive...).	You want Claude to act on data living in another app.

The single biggest jump in usefulness comes from **learning which face to reach for**. A report belongs in Cowork. A script belongs in Code. A repeated weekly task belongs in a Project.

2.2 The five levels — a map of where you are going

The five levels — a map of where you are going



A useful way to picture your progress. This book takes you comfortably to Levels 1–2; Level One reaches 3–4.

2.3 The one idea to take away

Old habit: the cleverness is in the question. New reality: the cleverness is in the **context** — how much of your work, your standards, and your files Claude can see before it answers. A good context with a plain prompt beats a brilliant prompt with no context. Almost everything in this manual is really about giving Claude better context.

3. Getting set up from zero

Chapter 1 rushed you through setup to get a quick win. Here it is properly, once.

1. **Get an account and a paid plan.** Cowork and Code need a paid plan. Prices below are current as of August 2026 — they change, so check claude.com.

Plan	Rough price	Good for
Free	€0	Trying Chat only. No Cowork.
Pro	~\$20 / month	One person, normal daily use. Enough for Cowork.
Max	~\$100–200 / month	Heavy daily use, long sessions, fewer limits.
Team / Enterprise	Per seat	A company with shared admin controls.

2. **Download the desktop app** from claude.com/download (Windows or macOS). Cowork lives in the desktop app, not the website.
3. **Sign in** with the same account everywhere so your projects and history follow you.

Use an email you will keep — you cannot change it later

Sign up with a personal email you will always control, not a work or school address you could lose. Claude ties your account, memory and history to that email and does not let you change it afterwards. People who used a company address lose everything the day they leave. (Switching from another AI? Chapter 13 shows how to bring your memory across.)

4. **Pick a model.** This is the engine. In a normal chat you can switch model per message — but in a Cowork conversation the model locks once you send the first message, so choose it before you start.

Model	Personality	When
Opus (5.5)	The smartest, most careful.	Anything that matters or needs reasoning. Default for real work.
Sonnet (5)	Fast and capable.	Quick drafts, simpler tasks, when you want speed.
Haiku (4.5)	Cheapest and quickest.	Trivial, high-volume jobs.

The model locks when a Cowork task starts

In Chat you can switch model between messages. In a Cowork conversation the model is fixed the moment you send the first message — to use a different one you start a new conversation (and lose that chat's built-up context). So choose deliberately up front. And mind the trade-off: a more capable model (Opus) gives better answers but uses more of your token allowance than Sonnet or Haiku.

How much can it read at once?

The most capable models (Opus and Sonnet) can hold about 1,000,000 tokens at once (this is called the context window) — very roughly 555,000 words on the newest models (about 750,000 on older ones), since a token is a little more than half a word; the small fast one (Haiku) holds about 200,000. That is huge — but, as Chapter 12 explains, fuller is not always better.

Model names and numbers keep moving — trust the picker

Think in tiers, not version numbers: a small fast one (Haiku), a balanced one (Sonnet), the most capable (Opus), and an even higher flagship family (Fable). Anthropic releases new versions constantly — Fable 5.1, Opus 5.5 and Sonnet 5 all arrived in 2026 — so the exact numbers on this page will drift. Whatever the model picker shows you today is the truth. For real work, pick the most capable model you have.

5. **Turn on Extended Thinking and Web Search.** Extended Thinking makes Claude reason before answering. Web Search lets it check facts instead of guessing — the best cure for making things up. Leave both on for serious work.

The 60-second rule

In my experience, Opus + Extended Thinking + Web Search turned on makes the biggest difference to your results. It takes one minute to set up.

4. Your first session, in detail

Chapter 1 was the sprint; this is the slow walk, so you understand each move.

4.1 Exercise A — summarize a document and save it

Copy-paste prompt

Read FILENAME in this folder. Before you start, ask me any questions you need. Then write a clear one-page summary in plain language, save it as a Word file in this same folder, and show me the link.

1. Claude may ask a short question first — answer it. (That habit is Chapter 5, and it is the most useful one in the book.)
2. It reads the file, writes the summary, and saves it as a new Word file in your folder.
3. It shows a clickable file card. Open it — that is your finished deliverable.
4. Refine: *"Make it half the length and add three key takeaways at the top."* You did not start over — you steered.

4.2 A complete example, end to end

Real work is usually more than one step. Here is a small business example that strings several together — a supplier's price offer arrives as a PDF and you need to act on it.

1. **Summarise.** "Read the offer PDF in this folder and give me a half-page plain-language summary: what is offered, the price, the deadline, and anything unusual."
2. **Pull the numbers out.** "Now extract the line items into an Excel file in this folder — description, quantity, unit price, total — and add a total row."
3. **Draft the reply.** "Draft a short, polite reply asking for a 5% discount and a firm delivery date. Keep it under 120 words."
4. **Check before you send.** "Re-read the offer and confirm the numbers in the Excel match it. Tell me anything you are unsure about."

What this shows

One folder, one conversation, three real deliverables (a summary, a spreadsheet, a draft email) — and a verification step at the end. This is the shape of most useful work in Cowork.

4.3 Exercise B — let Claude tidy up (optional)

Feel the "agent" side (Claude acting on its own, step by step): *"Look at the loose files in this folder. Propose a tidy subfolder structure and show me the plan before moving anything."*

Claude explains its plan first and waits for your yes.

If you can't find the file it made

It is almost certainly still in Claude's temporary workspace and was never saved into your folder. Say "save the final file into my folder and show me the link." More in Chapters 10 and 11.

5. Two habits that beat any prompt

People spend hours hunting for the perfect "magic prompt." It is wasted effort. Two simple habits do more for your results than any template. Build these two and everything else gets easier.

5.1 Let Claude ask YOU the questions first

The instinct of every beginner is to write the perfect instruction. The faster path is the opposite: let Claude interview you. It knows what it is missing; you usually do not.

End your request with four words: **"Ask me questions first."** Claude asks a few short questions — often as clickable choices — then does the work. Better result, far less typing.

Copy-paste this

Before you start, ask me any questions you need to do this well — one at a time — and wait for my answers before doing the work.

Second half: for anything big, make Claude show its plan before it acts. "Show me your plan first and wait for my go-ahead" turns a risky one-shot into something you can steer.

5.2 Never trust a confident answer — make Claude check

Claude sounds exactly as sure when it is wrong as when it is right. Confidence is a setting, not a fact-check. Treat a smooth, certain-sounding answer as a draft, not a verdict.

1. **Keep Web Search on** so Claude verifies facts against live sources instead of guessing from memory.
2. **Ask it to audit itself:** "Check your own answer before I trust it. Verify the facts and tell me what you are unsure about."

One more trap: Claude is **agreeable** — if you push back it may fold even when it was right. Don't take agreement as proof; ask it to argue the other side.

The lesson behind this chapter

It is not because something is written down that it is correct — that applies to notes, web pages, and anything Claude itself produces. Always check important claims against a reliable source, such as Anthropic's official help pages. This manual was fact-checked that way.

6. Common myths about Claude

Beginners lose the most time to a handful of wrong beliefs. Here they are, with the reality.

The myth	The reality
A longer chat remembers more.	It gets worse as it fills up. Start a fresh chat for a new topic (Chapter 12).
If Claude agrees, you must be right.	Claude is agreeable. Make it argue against you before you believe it.
Confident means correct.	It sounds just as sure when it is wrong. Confidence is a setting, not a fact-check.
You need to master prompting.	You need four words: "ask me questions first." Claude interviews you.
A longer prompt is a better prompt.	A clear goal beats a 200-word masterpiece. Say the destination, not every step.
The website and the desktop app are the same.	The desktop app does far more — Cowork and Code live only there.
Claude is just a chatbot.	Chat is one face of several (Chapter 2). Most people never leave the lobby.
Skills are only for developers.	A skill is just a playbook Claude runs for itself. No code required.
Pasting your logins helps Claude help you.	Never paste a password. Connect apps through connectors instead (Chapter 13).
The first draft is the deliverable.	The first draft is yours to fix, not to ship.
You are too far behind to start.	Open Claude tonight and do the Chapter 1 exercise. That is how everyone starts.

7. Where your files live — OneDrive and the big beginner trap

This is the chapter that trips up almost everyone. Understand it and half of the mysterious problems disappear.

7.1 Three different places

Where your files live when you work in Cowork



The rule that fixes everything: **a file only really exists for you once it is saved into the folder you selected.** If Claude says it made something but you cannot find it, it is almost always still in the scratchpad (Claude's private draft area) or the sandbox (a temporary, closed-off computer Claude works in) and needs saving to your folder.

7.2 OneDrive — friend and trap

If you keep your work in OneDrive (for example `C:\Users\you\OneDrive\Documents\...`), it backs everything up and syncs across devices — but it has two behaviours every beginner must know.

- **"Files On-Demand."** OneDrive can show a file while the real content lives only in the cloud (a cloud icon, not a green check). A tool may then read an empty placeholder. Fix: right-click the folder, choose "Always keep on this device."
- **Sync while writing.** If OneDrive syncs a file at the exact moment Claude is writing it, you can end up with a half-written file — one cause of the truncation problem in Chapter 11.

The OneDrive trap

Two machines editing the same OneDrive file, or OneDrive syncing mid-write, can create "conflict copies" or a file cut off halfway.

Habits that avoid it: keep the folder set to "Always keep on this device"; let a write finish before opening the file; and after a long file is created, open it once to confirm the end is really there.

7.3 Keep the folder simple

A messy folder makes Claude slower and more likely to read the wrong thing. The recommended shape for any general workspace is just three subfolders, explained in Chapter 8.

8. Cowork, explained simply

Cowork is the most important shift in how Claude is used. It is a desktop app that points at a folder on your computer and treats it as its workspace. It reads your files, writes new ones, runs code in the background, and asks you questions before doing anything important.

If Chat is a conversation, Cowork is more like an **employee**. You describe the outcome, it reads the relevant files, asks for clarifications through a clickable form, and comes back with a finished deliverable. As of mid-2026 it is generally available on all paid plans (Pro, Max, Team, and Enterprise) through the desktop app for Windows and macOS, which now shares one home with Chat — your Projects sit in the left sidebar, so if one seems to vanish after an update, look there or use search.

Scope its access — it can delete files

Cowork can read, edit, and permanently delete files in the folder you give it. Always point it at a specific project folder, never at your whole Documents or your entire drive.

8.1 The folder structure that saves hours

```
Claude Cowork/  ABOUT ME/      -> who you are, how you work, what to avoid
OUTPUTS/       -> finished deliverables (Claude organises by project)  TEMPLATES/
-> reusable structures Claude saves on request
```

- **ABOUT ME** is read at the start of every session. Keep these files short.
- **OUTPUTS** is where finished work lands. Cowork does not read it unless you point at a file — which keeps things fast.
- **TEMPLATES** fills itself: say "save this as a template" and it keeps the structure, not the content.

8.2 The first thing to type in a new project

Copy-paste this into any new Cowork project

I just created this project. Read every file in the folder. Then tell me, in plain language, what is here, what you think I use it for, and what instructions you will follow. If anything is unclear, ask me with a multiple-choice question before doing real work.

8.3 Connectors, scheduling, and working from your phone

- **Connectors** link Cowork to your live tools (Gmail, Outlook, Calendar, Drive, Notion, Slack). Connect only what a task needs.
- **Scheduled tasks** let a saved routine run on its own — a digest every weekday at 7am (your computer has to be awake).
- **Dispatch** lets you hand Claude a task from your phone; the desktop app does the work and you open the result later. Beta on Pro and Max; needs both the desktop and mobile apps.

8.4 Cowork vs Code — which one?

Short answer: **if the thing you want is a document, stay in Cowork. If the thing you want is software, use Code.** Cowork can still run code in its sandbox when a task needs it. Truncation is handled by good saving habits (Chapter 11), not by switching products.

9. Mistakes to avoid in Cowork

Seven habits that separate a smooth setup from a frustrating one, with the fix for each.

The mistake	Do this instead
Using Chat for real file work.	Point Cowork at the folder and let it produce the file. Chat is for thinking, Cowork for doing.
No "about me" file, so output sounds generic.	Put a short about-me and a list of words/phrases you never use in ABOUT ME.
Vague global instructions ("be helpful").	Be specific: where to save, what tone, when to ask, what never to do.
Opus for absolutely everything.	Use Opus where it matters; switch to Sonnet for short, simple tasks to save your allowance.
No skills installed.	Turn a task you do often into a reusable skill — one skill can replace twenty prompts.
Giving access to too much of your computer.	Scope Cowork to one project folder. It can delete files; keep the blast radius small.
No scheduled tasks.	Let recurring work (briefings, digests) run automatically while you sleep.

10. The sandbox and your security perimeter

When Claude needs to run something — convert a file, crunch numbers, build a spreadsheet — it does it inside a sandbox: a small, isolated Linux computer (Linux is an operating system, like Windows) that spins up just for your session.

10.1 What the sandbox is, and how much access you give

What Claude can see and do depends on the access you give

Sandbox (code)	Connected folder	Connectors	Computer use
Isolated code run. Can't touch the rest of your PC.	Read & write files. Can delete — only with your go-ahead.	Access to other apps (mail, calendar). Give minimal access.	Can act on your screen & visible apps. Maximum caution.

Golden rule: point Cowork at **ONE** project folder, and give only the access a task actually needs.

- **Isolated.** The sandbox cannot touch the rest of your PC — only the folder you connected and its own temporary space. That is what makes running code safe.
- **Temporary.** At the end of the session it is wiped, so files must be saved to your folder to survive.
- **A real computer.** It has Python and Node (programming languages) and common tools, so Claude can genuinely build and test, not just describe.

Your "security perimeter" is simply the sum of the access you grant — the folder you connect, the connectors you switch on, and whether you allow computer use (Claude moving your mouse and typing on your screen). Keep each as small as the task needs.


The disappearing-file trap

If Claude builds something and you cannot find it, it was probably created only in the sandbox or scratchpad and never copied to your folder. Say "save the final file into my folder and show me the link."

11. Truncation — the "trunking" problem, solved

Truncation is when a long file gets **silently cut off** — usually the end is missing — while Claude reports that it saved successfully. You only find out when you open the file and the last pages are gone. It is avoidable.

“Saved successfully” is not proof the file is complete



11.1 Why it happens (possible causes)

- A synced folder (like OneDrive) saves a half-written copy while the real content is still being written.
- A format converter (for example Markdown, a plain-text format with simple formatting, to Word) occasionally drops the tail of the document.
- A tool says "success" before the whole file has actually landed on disk.

11.2 How to prevent it

1. **Ask for files, not walls of text.** For anything long, have Claude write to a file rather than printing it all in the chat.
2. **Make Claude verify the end.** "Re-open the file and confirm the last section is really there."
3. **Convert long PDFs to Markdown before uploading.** Cleaner input, less truncated output.
4. **Keep OneDrive calm.** "Always keep on this device" and let writes finish (Chapter 7).

A reliable backup if a download button fails

Open the file in the Claude interface, select all, copy the raw text into a plain editor (e.g. Notepad++), and save it with the right filename.

12. Sessions, limits, and good habits

Claude is powerful but not unlimited. Knowing how sessions and limits work keeps you from hitting a wall mid-task.

12.1 The two limits

Limit	What it means
5-hour rolling limit	A cap on how much you can do within any 5-hour window. It refills as time passes.
7-day weekly limit	A larger cap across the week. Heavier plans (Max) have more headroom.

Timing

In my experience, the busy window in Europe is roughly 14:00–20:00. When you can, do heavy sessions in the morning — it is faster and you waste less of your allowance waiting.

12.2 The "smart zone" — why fuller is not better

Even though the context window is huge, Claude is sharpest when it is not crammed full. As a conversation fills with files and back-and-forth, answers slowly get duller. Keep each task small and start clean often rather than letting one chat swell to the brim.

12.3 How long to stay in one session

My rule of thumb: start a fresh chat when the conversation passes roughly 50 messages or you have uploaded several big files; when you change topic; or when Claude starts repeating itself. A good rhythm: one topic per session, an hour or two at most.

12.4 Session notes — never explain yourself twice

Because each new chat starts blank, keep a short session-notes file per project (for example `project_session_notes.md`). Start a new chat by uploading it and saying "Continue this project — read the notes above," and ask Claude to update it at the end. Chapter 13 explains how memory does the same thing automatically.

12.5 Habits that stretch your allowance

- **Keep instruction files short** (in my experience, under 500 words works best).
- **Batch your asks** — several related requests in one message.
- **Convert PDFs to Markdown** before uploading; they cost far fewer tokens.
- **Turn off connectors you are not using** so they don't quietly eat context.

13. Projects, Skills, and Memory

Three features that turn Claude from a clever stranger into a colleague who remembers.

13.1 Projects

A Project is a folder of chats that share instructions and files. Make one per recurring deliverable — not one giant project for everything. Upload only what that deliverable needs, and write five to eight lines of instructions describing what it is, who it is for, and the quality bar. Since mid-2026 your projects live in the left sidebar, shared between Chat and Cowork — pin the ones you use so they are easy to find.

13.2 Skills

A Skill is your standard way of doing one job, written so Claude runs it the same way every time. If you already have a checklist, you basically have a skill.

- **It is just a file.** A skill lives in a small text file (SKILL.md). Only a one-line description sits in Claude's memory until a task matches it — then the full instructions load.
- **Build it from real work.** "Turn what we just did into a reusable skill, and add a section of gotchas." Claude writes the file, not you.
- **You manage them** in the app, under Customize, then Skills. A skill only improves when its file is edited.

The easiest way to make a skill: record yourself doing it

New in Cowork (mid-2026): open the + menu in the desktop app and choose "Record a skill." Claude records your screen and your voice while you do the task once — you talk through your choices as you go — and turns it into a reusable skill. No writing required, which makes it the easiest on-ramp for a job that is simpler to show than to describe. Do not record passwords, secrets, or private information. Available on Pro, Max and Team.

13.3 Memory and session notes

Memory lets Claude carry facts between sessions. A MEMORY.md file plus small notes capture your setup and preferences, so a brand-new session already knows who you are — the automatic version of the session-notes habit.

Coming from ChatGPT? Import your memory in one minute

If you are switching from ChatGPT, Gemini or Grok, you do not have to start from zero. Go to claude.com/import-memory, copy the short prompt it gives you, paste that into your old assistant, then copy its answer back into Claude under Settings, Capabilities, Memory. Claude folds in what it learns (it can take a few hours to settle). Do this once when you switch, then keep using the session-notes habit above. On Pro and Max.

14. Safety and good practice

14.1 Classify your data before you paste

Sensitive data — classify before you paste

Chat, Cowork, and Code are about workflow, not confidentiality: all three send your content to Anthropic's servers, so "it runs on my computer" does NOT make sensitive data safe. Before pasting, classify it — green (public/harmless), orange (internal/operational), red (personal data, legal or warranty disputes, contracts, accident files). Keep red out, anonymise it, or get approval first. The one-page rule is reproduced in the Annex.

14.2 The rest of the checklist

- **Work in a dedicated folder.** Point Cowork or Code at a project folder, never your whole Documents.
- **Be careful with "skip permissions" (letting Claude act without asking you first).** Only turn it off where the files could be replaced.
- **Never paste secrets.** API keys (secret codes that let a program use a service) and passwords belong in a local settings file, not a chat.
- **Keep a backup.** OneDrive version history and timestamped ZIP backups mean a bad edit is never fatal.
- **Verify anything that matters.** For numbers, legal or money decisions, check the source — and use Web Search so Claude checks too.

Work you build on the job

Anything built on company time or equipment can raise ownership questions. Agree up front who owns it — and if it is meant to be personal, keep it on personal equipment and time.

14.3 Can I keep my data local?

Short answer: no — but you can limit its use

Chat, Cowork and Code all send your content to Anthropic's servers to be processed. Claude does not run on your computer, and no setting makes your data stay local. What you CAN do: stop your conversations being used to train the models, which also shortens how long they are kept. For truly local, offline AI you would need a different, self-hosted model — which is not Claude.

How to opt out of training

1. Open Settings, then Privacy. 2. Turn off the model-training toggle (often labelled 'help improve Claude' or 'use my data to improve Claude'). 3. It applies to new and resumed conversations going forward. Team, Enterprise and API accounts do not train on your data by default. Note: content flagged for safety review may still be used, regardless of this setting.

Annex — AI Data Rules (one page)

Keep this rule in front of anyone using Claude for work. The key idea: Chat, Cowork and Code are about workflow, not confidentiality — all three send your content to Anthropic's servers, so classify the data first.

Light	Means	Examples	What to do
GREEN	Public / harmless	Public info, marketing text, general questions.	Use freely.
ORANGE	Internal / operational	Internal procedures, non-personal company data.	Approved tools; avoid detail.
RED	Sensitive	Personal data, legal/warranty disputes, contracts.	Anonymise, get approval, or keep out.

1. Confidential material only in approved tools and accounts.
2. A human always reviews before anything is used or sent.

15. Quick-start cheat sheet

Your first 30 minutes

1. Install the desktop app from claude.com/download and sign in.
2. Set the model to Opus; turn on Extended Thinking and Web Search.
3. Create a working folder with ABOUT ME / OUTPUTS / TEMPLATES inside it.
4. Open Cowork, point it at that folder.
5. Do the Chapter 1 exercise: summarize a document and save it.
6. Paste the "read every file" prompt from Chapter 8.2 to brief Claude on the folder.

Prompts worth keeping

Let Claude ask first

Before you start, ask me any questions you need — one at a time — and wait for my answers before doing the work.

Make it check itself

Audit your own answer before I trust it. Verify the facts with web search and tell me what you are unsure about.

Safe long file

Write this to a file in my folder, then re-open it and confirm the last section is complete before you tell me it's done.

16. Plain-language glossary

AI words

Term	In plain words
Model	The engine doing the thinking (Opus, Sonnet, Haiku). Smarter engines use more of your allowance.
Prompt	What you type. The kick-off, not the whole job.
Context	Everything Claude can see while answering. The real source of quality.
Token	A chunk of text, roughly half to three-quarters of a word, depending on the model. Limits and costs are counted in tokens.
Context window	How much text Claude can hold at once (about 1M tokens on Opus/Sonnet, 200K on Haiku).
Extended Thinking	A mode where Claude reasons before answering.
Web Search	Lets Claude check facts online instead of guessing.
Agent	Claude working toward a goal on its own, step by step.

Where your files live

Term	In plain words
Cowork	The desktop agent that works inside a folder and produces finished files.
Sandbox	A temporary, isolated Linux computer Claude uses to run code safely.
Scratchpad	Claude's private draft area, before a file is saved into your folder.
Truncation	A file getting silently cut off, usually at the end. See Chapter 11.
Files On-Demand	A OneDrive setting where the real file stays in the cloud until you open it.
Markdown	A plain-text format with simple formatting. Cheap in tokens, easy for Claude to read.

Connecting, remembering and staying safe

Term	In plain words
Project	A folder of chats sharing instructions and files.
Skill	A reusable instruction file (SKILL.md) that teaches Claude one job.
Memory / MEMORY.md	Facts Claude keeps between sessions so it remembers your setup.
Connector	A bridge that lets Claude reach another app, like Gmail or Outlook.
MCP	The shared standard that connectors use to talk to Claude.
Dispatch	Starting a Cowork task from your phone; the desktop does the work.
Code	Claude that writes and runs real software.

Term	In plain words
API key	A secret code that lets a program use a service. Never paste it in a chat.
Skip permissions	A mode where Claude acts without asking first. Use it only on files you can replace.
Computer use	Claude moving your mouse and typing on your screen. Give it the least access you can.

Where to learn more — video tutorials

If you prefer watching, these beginner-friendly videos (search for the titles on YouTube) are a good next step. They are in English, where the most information is.

- FULL Claude Cowork Tutorial for Beginners — set up a folder first, and watch Cowork ask questions before it acts.
- Learn 97% of Claude in 16 Minutes — a fast tour of the whole product.
- Every Level of Claude Explained — the five-level map from Chapter 2.

A note on these videos

Several name older models and quote out-of-date numbers. For model names, trust the model picker in the app over any video or book.

What's next — Level One

This book gets you comfortable at Levels 1–2 of the map in Chapter 2: set up Claude, run a task in Cowork, find your files, avoid truncation, manage sessions. Level One moves you to Levels 3–4 — real prompting, Projects in depth, your first Skill, connectors, a gentle step into Code and Git, and your first small agent.

In one line: Level Zero is learning to drive; Level One is planning routes and building the car. Finish the Chapter 1 exercise, use Claude for a week, then move up.

By Hans Tandt