

Choosing the Right Claude Model

A field guide: what to use, why it costs what it costs, and how the limits actually work

Version 1.3 · 28 September 2026



Contents

1. Purpose of this guide	3
2. The quick rule, and what it leaves out	3
3. The current model lineup	4
4. Quick decision guide	4
5. Why switching models is worth doing deliberately	4
6. What is a token?	5
7. Token and compute costs	5
8. Plans: private vs. corporate accounts	6
8.1 Private (individual) plans	6
8.2 Corporate (business) plans	6
9. Session limits, weekly limits, and where the reset clock comes from	7
10. Models change over time, and newer ones can burn through a quota fast	7
11. Self-running agents: why the monthly bill can move faster than expected	8
12. What else is worth knowing	8

NOTE — Created 5 September 2026 - last updated 28 September 2026

Model names, prices and limits change often. Figures here were checked against Anthropic's official documentation as of the last-updated date; verify at claude.com/pricing and platform.claude.com/docs before relying on exact numbers.

1. Purpose of this guide

This guide helps you pick the right Claude model and explains what sits behind that choice: why it is worth taking seriously, what a token actually is, how tokens turn into cost, how the private and company Claude plans differ, and how the usage limits and their resets actually work. It closes with a short list of things a quick decision chart tends to leave out.

2. The quick rule, and what it leaves out

A popular way to picture the choice is a simple decision tree (one widely shared version is an infographic by Dr. Joerg Storm on LinkedIn; it is not an Anthropic publication): for a quick answer, pick Haiku for speed or Sonnet for better quality; for anything harder, pick Opus for a single hard reasoning task, or Fable for work that spans many steps and tools.

That logic holds up well against Anthropic's own guidance as of September 2026. The four model families exist, and the ordering matches how Anthropic frames the choice: Haiku fastest and cheapest, Sonnet the best mix of speed and intelligence, Opus for serious work (Anthropic now suggests starting with Opus for most workloads), and Fable for the most demanding reasoning and long, tool-heavy agentic work. A few things are worth adding, though.

- Version numbers are not shown, and they matter. As of this writing the current releases are Haiku 4.5, Sonnet 5, Opus 5.5 and Fable 5.1. Each new version changes what good enough means, so use Sonnet is really shorthand for use whichever Sonnet is current.
- Fable is not available on every plan, and it is capped. It is included from Pro upwards (Pro, Max, Team and Enterprise), but on Pro and Max it can use at most about half of your weekly limits. The Free plan offers Haiku and Sonnet only.
- No cost information. The chart is purely about speed versus quality trade-offs; it says nothing about price, which differs by up to 10 times between the cheapest and most expensive model.
- Extended, or adaptive, thinking is a separate lever from model choice. Turning up how much a model reasons before answering raises both quality and cost independently of which of the four you picked.

The practical rule behind any such chart is sound: use the lightest model that reliably does the job, and move up only when the task itself gets more complex, ambiguous, or higher-stakes, not as a default habit.

3. The current model lineup

Model	Best used for	Context window	Relative cost
Haiku 4.5	Quick chat, summaries, first drafts, high-volume or low-cost tasks	200K tokens	Lowest
Sonnet 5	Fast everyday work: writing, analysis, planning, most business tasks and coding	1M tokens	Low to medium
Opus 5.5	Serious work, hard reasoning, nuanced trade-offs and long agentic coding; Anthropic's suggested starting point	1M tokens	Medium
Fable 5.1	Long, multi-step, tool-heavy agentic work and orchestration; reserve for the hardest 10%	1M tokens	Highest

NOTE — Context windows differ by model

Haiku keeps the older 200K context window; Sonnet, Opus and Fable currently offer 1M tokens. Check platform.claude.com/docs/en/models for current figures.

4. Quick decision guide

Situation	Use	Why
A quick fact, summary, or first draft	Haiku	Fastest and cheapest; more than good enough for low-stakes text
Everyday writing, analysis, planning, most coding	Sonnet	Best balance of quality and cost; the sensible default for a team
The task is hard, ambiguous, or the answer really matters	Opus	Worth the extra cost when getting it right matters more than getting it fast
The task spans many steps, tools, or a long research chain	Fable	Built for multi-step, tool-heavy work; reach for it only where the job is genuinely that complex, and where your plan includes it

5. Why switching models is worth doing deliberately

It is tempting to always pick the biggest model to be safe. In practice the choice affects four things at once.

- Speed: lighter models answer faster, which matters for anything interactive or high-volume.
- Quality and depth: heavier models reason more carefully and handle ambiguity, nuance and long chains of steps better, but most everyday tasks do not need that depth.
- Cost multiplier: Fable's output tokens cost roughly 10 times what Haiku's do, and Opus roughly twice what Sonnet does. Running everything through the top model is not just slower at scale; it is materially more expensive.
- Context capacity: Haiku's 200K-token window is noticeably smaller than the 1M available on the other three, which matters for very long documents, codebases, or conversations.

The practical rule captures this well: use the lightest model that reliably does the job, and move up only when the task itself demands it.

6. What is a token?

A token is the basic chunk of text a model reads and writes; not quite a word, not quite a character. Claude's tokenizer breaks text into sub-word pieces, so as a rough rule of thumb for English: on the newest models 1 million tokens is roughly 555,000 words, so one token is a little more than half a word (older models fitted about 750,000 words in the same space). An unusual family name might be one token or split into two or three pieces; uncommon words, source code, and non-English languages typically tokenize less efficiently, meaning more tokens for the same amount of meaning.

Images and PDF pages are also converted into a token-equivalent cost based roughly on their size and resolution, even though that number is never shown directly in a chat. Tokens, not words, characters, or messages, are what actually get counted and billed, both for what you send (input) and what the model produces (output).

7. Token and compute costs

Model	Input per million tokens	Output per million tokens
Haiku 4.5	\$1	\$5
Sonnet 5	\$2	\$10
Opus 5.5	\$4	\$20
Fable 5.1	\$10	\$50

NOTE — Subscriptions do not bill per token directly

These are official developer or API list prices. A subscription plan (section 8) charges a flat monthly fee for a usage allowance, but the same cost differences between models drive how large that allowance feels in practice.

Two things tend to surprise people who are new to this.

- Output costs more than input, model for model, usually about 5 times. Reading text is cheap and can happen in parallel; generating it is slower and more compute-heavy, so every model charges more per token for what it writes than for what you send it.
- Thinking tokens are invisible but billed. Opus and Fable especially, and Sonnet to a lesser degree, can reason step by step before answering. That reasoning is made of output tokens too, even though the chat only shows the final answer.
- Two levers cut the bill without changing the model: prompt caching, which reuses previously-sent context at a steep discount, and the batch API, which trades slower turnaround for a substantial discount on large offline jobs.

8. Plans: private vs. corporate accounts

8.1 Private (individual) plans

Plan	Indicative price	Models included	What it adds
Free	\$0	Sonnet, Haiku	Chat on web, desktop and mobile; web search, files and code; memory; connected apps; a modest session allowance
Pro	About \$20/mo (\$17/mo billed annually)	+ Opus, Fable (capped)	Everything in Free, plus Claude Code, Claude Design, Projects, Research; at least 5 times the usage of Free per session
Max 5x	\$100/mo	Same as Pro, Fable capped	Everything in Pro at 5 times the usage; earlier access to new features
Max 20x	\$200/mo	Same as Pro, Fable capped	Everything in Pro at 20 times the usage; earlier access to new features

NOTE — Fable is capped on Pro and Max plans

Fable use is capped at roughly half of the weekly quota; it cannot consume the entire weekly budget on its own.

8.2 Corporate (business) plans

Plan	Indicative price	What it adds
Team	From about \$25/mo per standard seat (\$20/mo annually); about \$125/mo per premium seat (\$100/mo annually)	Central billing, admin controls, SSO, a shared workspace; all models available
Enterprise	From about \$20/seat/month plus usage-based charges, billed annually	Everything in Team, plus SCIM, audit logs, a compliance API, custom data retention, HIPAA-readiness, role-based access

The practical difference for a company: a Team or Enterprise seat gives every colleague their own quota, with the organization able to see aggregate usage and control who has access to which models and features. Individual Free, Pro or Max plans are personal, self-service, and not centrally managed.

9. Session limits, weekly limits, and where the reset clock comes from

Paid plans meter usage two ways at once.

- A 5-hour session window. It starts the moment you send your first message, and every message sent in the following 5 hours counts against the same window, whether it is one long conversation or several separate chats.
- A weekly limit on top of that: a larger quota that renews on a roughly weekly cycle, generally split into an overall pool for all models plus a smaller pool specifically for the heaviest model on your plan.

Where the exact day and time of your personal reset comes from: Anthropic does not publish one fixed global clock that resets everyone's weekly limit at the same moment. It is specific to each account, and the exact countdown is shown under Settings, then Usage. In practice it behaves like a rolling window anchored to when your current weekly cycle started, rather than a fixed calendar moment such as every Monday at 00:00 UTC. Some users have reported it landing slightly earlier than a full 7 times 24 hours later, so treat the in-app countdown as the source of truth for your own account rather than something you can work out from first principles.

It is also worth separating three different things that all get called a session.

- The 5-hour usage-limit session above: a billing and quota concept.
- A single conversation or task session, which can run far longer than 5 hours of wall-clock time while only drawing down usage for the messages actually sent within it.
- A model's context window: how much text, input and output combined, fits into one conversation before older content has to be summarized or dropped. This is a size limit, not a time-based or renewing one.

10. Models change over time, and newer ones can burn through a quota fast

None of the four models in this guide is permanent. Anthropic retires older versions on a rolling basis and replaces them with new ones, the same way Opus 5.5 replaced Opus 5, and Sonnet 5 replaced Sonnet 4.6 before it. Haiku 4.5, for example, is only guaranteed until 15 October 2026. Once a version is retired it eventually stops being callable at all, even though it may still be named in an older document like this one. Treat every table above as a snapshot, not a fixed reference; it is worth rechecking platform.claude.com/docs/en/models every few months.

TRAP — A newer, more powerful model can burn a weekly quota fast

It is genuinely common to try the newest top-tier model (Fable, currently) on an open-ended multi-step task and watch a large part of a week's quota disappear in well under an hour, sometimes within minutes, because it costs more per token and keeps working through many tool calls without pausing.

The practical takeaway: the first time you or your team try a newer or more powerful model on an unattended or long-running task, keep Settings, then Usage open and watch it for the first few minutes rather than assuming it will behave like the model you are used to.

11. Self-running agents: why the monthly bill can move faster than expected

Everything so far assumes a person is sending each message. As soon as an agent starts triggering itself, such as a scheduled task or a recurring workflow, cost stops being about a single conversation and becomes a multiplication of three things that can each move independently: how often it runs, how many tokens it uses per run, and which model is doing the work.

Each factor on its own can be a large multiplier. Going from running once a day to every 15 minutes is a 96 times increase in frequency alone. Moving the same recurring task from Haiku to Fable is up to a 10 times increase in price per token alone. Because these multiply rather than add, changing both at once does not add the two effects; it multiplies them, so a small, cheap-looking experiment can turn into a bill hundreds of times larger without anyone deciding that on purpose.

A fourth factor makes this easy to underestimate further: tokens per run is rarely fixed. An agent that calls tools, reads files, or resends a growing history on every run can use far more tokens per run than a simple one-shot prompt, and that number tends to grow over the agent's lifetime rather than staying constant.

TRAP — Two different failure modes

An agent on a subscription seat draws from the same 5-hour and weekly quota as interactive use, and can lock the team out for the rest of the week. An agent on the API has no such ceiling, so the risk flips to an open-ended invoice instead.

TIP — Before letting anything run on a schedule

Use the lightest model that reliably does the recurring task, turn on prompt caching for any large fixed context, and check real per-run token usage early.

12. What else is worth knowing

- Context window and usage limit are two different ceilings. One caps how much fits in a single conversation; the other caps how much can be sent over time.
- API or developer usage is not capped by the 5-hour or weekly limits at all; it is pay-as-you-go per token with no ceiling, so the risk there is a large bill, not a lockout.
- Rate limits, meaning requests per minute and tokens per minute, are a separate constraint on the API side, on top of per-token cost.
- Vision and file inputs such as images, PDFs and screenshots consume tokens too, sometimes a lot of them for a large image or a long PDF.
- Data handling differs by plan. Enterprise offers custom data retention and HIPAA-readiness that Free, Pro and Max do not.
- Team and Enterprise admins get an organization-wide usage dashboard to see spend and activity, versus an individual plan where only that person sees their own usage.

Figures were checked against Anthropic's official pages on the date on the cover. Before choosing a plan, check claude.com/pricing: names, prices and limits change often.